

# Lab 4

---

**Members: Fangwen Liao (3439869) | Yijin Wang (3476217)**

## Exercise 1

(a) According to man page of strace, it is used to intercept and record system calls made by a process. Run the following command in terminal and record the result in fork.log:

```
fangwenliao@debian:~/Downloads/OS_Lab2/fork$ strace -o fork.log ./fork
Main process PID: 1047
Child PID: 1048
Child PID: 1049
Child PID: 1050
Child PID: 1051
Child PID: 1053
Child PID: 1054
Child PID: 1052
Press ENTER key to Continue

Process 1051 ended

Process 1048 ended

Process 1049 ended

Process 1050 ended

Process 1053 ended

Process 1054 ended

Process 1052 ended

Process 1047 ended
```

The content in the log file is below:

```
1 execve("./fork", [ "./fork" ], [ /* 20 vars */ ]) = 0
2 brk(NULL)                                = 0x62c000
3 access("/etc/ld.so.nohwcap", F_OK)       = -1 ENOENT (No such file or
directory)
4 mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0xb77a2000
5 access("/etc/ld.so.preload", R_OK)      = -1 ENOENT (No such file or
directory)
6 open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
7 fstat64(3, {st_mode=S_IFREG|0644, st_size=99587, ...}) = 0
```

```

8 mmap2(NULL, 99587, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb7789000
9 close(3) = 0
10 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or
directory)
11 open("/lib/i386-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
12 read(3,
"\177ELF\1\1\1\3\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\204\1\0004\0\0\0"... , 512) =
512
13 fstat64(3, {st_mode=S_IFREG|0755, st_size=1791908, ...}) = 0
14 mmap2(NULL, 1800700, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) =
0xb75d1000
15 mprotect(0xb7782000, 4096, PROT_NONE) = 0
16 mmap2(0xb7783000, 12288, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1b1000) = 0xb7783000
17 mmap2(0xb7786000, 10748, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7786000
18 close(3) = 0
19 set_thread_area({entry_number:-1, base_addr:0xb77a3100, limit:1048575,
seg_32bit:1, contents:0, read_exec_only:0, limit_in_pages :1,
seg_not_present:0, useable:1}) = 0 (entry_number:6)
20 mprotect(0xb7783000, 8192, PROT_READ) = 0
21 mprotect(0x4aa000, 4096, PROT_READ) = 0
22 mprotect(0xb77cb000, 4096, PROT_READ) = 0
23 munmap(0xb7789000, 99587) = 0
24 getpid() = 1047
25 fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
26 brk(NULL) = 0x62c000
27 brk(0x64d000) = 0x64d000
28 write(1, "Main process PID: 1047\n", 23) = 23
29 clone(child_stack=NULL, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
child_tidptr=0xb77a3168) = 1048
30 clone(child_stack=NULL, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
child_tidptr=0xb77a3168) = 1051
31 clone(child_stack=NULL, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
child_tidptr=0xb77a3168) = 1054
32 getpid() = 1047
33 nanosleep({tv_sec=1, tv_nsec=0}, 0xbfca0ae8) = 0
34 write(1, "Press ENTER key to Continue\n", 28) = 28
35 nanosleep({tv_sec=1, tv_nsec=0}, 0xbfca0ae8) = 0
36 fstat64(0, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
37 read(0, 0x62c410, 1024) = ? ERESTARTSYS (To be restarted if
SA_RESTART is set)
38 --- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_EXITED, si_pid=1051, si_uid=1000,
si_status=0, si_utime=0, si_stime=0} ---
39 read(0, 0x62c410, 1024) = ? ERESTARTSYS (To be restarted if
SA_RESTART is set)
40 --- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_EXITED, si_pid=1048, si_uid=1000,
si_status=0, si_utime=0, si_stime=0} ---
41 read(0, 0x62c410, 1024) = ? ERESTARTSYS (To be restarted if
SA_RESTART is set)
42 --- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_EXITED, si_pid=1054, si_uid=1000,
si_status=0, si_utime=0, si_stime=0} ---
43 read(0, "\n", 1024) = 1
44 write(1, "Process 1047 ended\n", 19) = 19

```

```
45 exit_group(0)                = ?
46 +++ exited with 0 +++
```

In fork.c line 9 calls clone(), the corresponding line in log is line 29: clone(child\_stack=NULL, flags=CLONE\_CHILD\_CLEARTID|CLONE\_CHILD\_SETTID|SIGCHLD, child\_tidptr=0xb77a3 168) = 1048

In fork.c line 16 calls getpid() and write(), this is called by child processes, thus the trace are not found in this log file (to increase readability of log file, the -f option is not used, however when switched on those calls will be recorded).

In fork.c line 18 calls nanosleep(), the corresponding line in log is line 33: nanosleep({tv\_sec=1, tv\_nsec=0}, 0xbfca0ae8) = 0

In fork.c line 22 also calls nanosleep(), the corresponding line in log is line 35: nanosleep({tv\_sec=1, tv\_nsec=0}, 0xbfca0ae8) = 0

In fork.c line 24 calls read(), and since it is getchar() so the corresponding line in log is line 43, not the ones before: read(0, "\n", 1024)

(b) The clone() in line 9: According to man page of fork(2), at the time of fork the child has the same memory content as parent, so they execute the same stack, so the child\_stack parameter should be NULL, according to man page of clone(2), CLONE\_CHILD\_CLEARTID clear the child thread ID when child exits, CLONE\_CHILD\_SETTID store the child tid, SIGCHLD sent this signal to parent when child dies, child\_tidptr stores the child tid in parents memory. Return value is the child's pid if success.

The getpid() in line 16: According to man page of getpid(2), it returns the pid of the calling process.

The write() in line 16: to get the log, another strace with -f is runned and one write() call of a child is picked out as below, 1277 write(1, "Child PID: 1277\n", 16 <unfinished ...> 1277 <... write resumed> ) = 16 According to man page of write(2), the first argument is file descriptor which is a reference number of the operating file, in this case the fd is 1, [which is the one for standard output](#), the second argument is the start pointer of the content to write, and the third argument is the count of how many bytes to write from the start pointer. The return value is the number of bytes written, which is the third argument, when success.

The nanosleep() in line 18 and 22: According to man page in nanosleep(2), this call is used to suspend the process for a time specified in the argument, which is a structure consists of second and nanosecond, which provides high precision. Return value is 0 when successful.

The read in line 24: According to man page of read(2), the first argument is also a file descriptor, in this case is 0, [which is the one for standard input](#), the second argument is the content to read and the third one is the count number of bytes to read. Return value is the number of byte read, \n count as one byte.

## Exercise 2

(a)

```
#include <fcntl.h>
#include <unistd.h>
```

```
int main() {
    int fd = open("./output.txt", O_RDWR|O_CREAT|O_APPEND, S_IRWXU);
    write(fd, "Write something!\n", 17);
    close(fd);
}
```

First included header is `fcntl.h`, which [define](#) the `O_APPEND` `O_RDWR` and `O_CREAT` needed in `open()`, the second included header is `unistd.h`, which [defines](#) `close()` and `write()`.

The first argument of `open()` is the path of the file to open, the second are flags: one of the access mode flag must included, in our case is read/write. `O_RDWR`, `O_APPEND` is used to write content at the end of the file, `O_CREAT` is used to create the file when the file doesn't exist, when using this flag a mode must be specified, in our case the `S_IRWXU` is used, which grants user read, write and execution access. A `int` which indicates the file descriptor is returned.

`write()` is discussed before.

`close()` according to man page of `close(2)` takes one argument which is the file descriptor, return 0 when success. The same `strace` command is used as before and the log file is below:

```
1 execve("./a.out", ["/a.out"], [/* 20 vars */]) = 0
2 brk(NULL)                                = 0x18c7000
3 access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or
directory)
4 mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0xb7771000
5 access("/etc/ld.so.preload", R_OK)        = -1 ENOENT (No such file or
directory)
6 open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
7 fstat64(3, {st_mode=S_IFREG|0644, st_size=99587, ...}) = 0
8 mmap2(NULL, 99587, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb7758000
9 close(3)                                  = 0
10 access("/etc/ld.so.nohwcap", F_OK)       = -1 ENOENT (No such file or
directory)
11 open("/lib/i386-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
12 read(3,
"\177ELF\1\1\1\3\0\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\204\1\0004\0\0\0"...
, 512) = 512
13 fstat64(3, {st_mode=S_IFREG|0755, st_size=1791908, ...}) = 0
14 mmap2(NULL, 1800700, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) =
0xb75a0000
15 mprotect(0xb7751000, 4096, PROT_NONE)    = 0
16 mmap2(0xb7752000, 12288, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1b1000) = 0xb7752000
17 mmap2(0xb7755000, 10748, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7755000
18 close(3)                                  = 0
19 set_thread_area({entry_number:-1, base_addr:0xb7772100, limit:1048575,
seg_32bit:1, contents:0, read_exec_only:0, limit_in_pages :1,
seg_not_present:0, useable:1}) = 0 (entry_number:6)
20 mprotect(0xb7752000, 8192, PROT_READ)    = 0
```

```

21 mprotect(0x46b000, 4096, PROT_READ)      = 0
22 mprotect(0xb779a000, 4096, PROT_READ)     = 0
23 munmap(0xb7758000, 99587)                 = 0
24 open("./output.txt", O_RDWR|O_CREAT|O_APPEND, 0700) = 3
25 write(3, "Write something!\n", 17)         = 17
26 close(3)                                  = 0
27 exit_group(0)                             = ?
28 +++ exited with 0 +++

```

As we can see above, from line 24 to line 26 the system calls are successful.

(b) The argument in `open()` are explained before, it returns `fd = 3`, which is after the `fd of standard error = 2`. The argument and return value of `write()` is explained before. The argument and return value of `close()` is explained before.

(c) User mode can't access directly to hardware devices, only kernel mode can, thus system call is needed. System calls add a layer of protection to the system, so the user can't easily crash the system.

## Exercise 3

(a) Run the following command in terminal:

```
fangwenliao@debian:~$ watch -n.1 "cat /proc/interrupts"
```

```
Every 0.1s: cat /proc/interrupts
debian: Tue Dec  7 23:03:48 2021
```

	CPU0	CPU1	CPU2	CPU3			
0:	35	0	0	1	IO-APIC	2-edge	timer
1:	1	0	0	9	IO-APIC	1-edge	i8042
8:	0	0	0	0	IO-APIC	8-edge	rtc0
9:	0	0	0	2	IO-APIC	9-fasteoi	acpi
12:	0	0	0	162	IO-APIC	12-edge	i8042
14:	0	0	0	0	IO-APIC	14-edge	ata_piix
15:	0	0	0	469	IO-APIC	15-edge	ata_piix
16:	0	3993	0	62	IO-APIC	16-fasteoi	enp0s8
18:	0	0	0	0	IO-APIC	18-fasteoi	vmwgfx
19:	0	4	173	7	IO-APIC	19-fasteoi	enp0s3
20:	0	0	0	193	IO-APIC	20-fasteoi	vboxguest
21:	0	0	0	8722	IO-APIC	21-fasteoi	
ahci[0000:00:0d.0], snd_intel8x0							
22:	0	0	0	27	IO-APIC	22-fasteoi	
ohci_hcd:usb1							
NMI:	0	0	0	0	Non-maskable interrupts		
LOC:	14967	10315	6945	8643	Local timer interrupts		
SPU:	0	0	0	0	Spurious interrupts		
PMI:	0	0	0	0	Performance monitoring		
interrupts							
IWI:	0	0	0	0	IRQ work interrupts		

RTR:	0	0	0	0	APIC ICR read retries
RES:	5885	3888	3884	3227	Rescheduling interrupts
CAL:	2486	2380	3062	1133	Function call interrupts
TLB:	456	464	343	426	TLB shootdowns
TRM:	0	0	0	0	Thermal event interrupts
THR:	0	0	0	0	Threshold APIC interrupts
DFR:	0	0	0	0	Deferred Error APIC interrupts
MCE:	0	0	0	0	Machine check exceptions
MCP:	2	2	2	2	Machine check polls
ERR:	0				
MIS:	0				
PIN:	0	0	0	0	Posted-interrupt notification event
PIW:	0	0	0	0	Posted-interrupt wakeup event

**Each column has its meanings:** The first one is IRQ number, the next four are interrupt numbers of each CPU, the next is the type of interrupts and the last is the device that caused the interrupt.

### IRQ\_0 timer

(a) A software clock, update time and date, update time passed after system start. When a process is running, the system timer record its running time, and when it exceed its allocated time, the timer will send a interrupt to [preempt it](#).

### IRQ\_1 i8042

(a) Keyboard, [Handels keyboard input](#). (b) When press and release a key, the system will call two interrupt to handel the input content.

### IRQ\_8 rtc0

(a) Real-time clock, Another hardware clock, (b) according to man page of `rtc(4)`, it record the wall clock time, its has own backup power when the machine is turned off, it is not system lock. It will generate interrupts when a previously set alarm time is reached.

### IRQ 9 acpi

(a) ACPI stands for advanced Configuration and Power Interface. (b) It [provide a open standard for OS to perform power management](#). It defines a hardware abstraction interface between device firmware, computer hardware and OS. It raise interrupts when a General purpose event happens such as [plug the AC adapter or close the lid of laptop](#).

### IRQ 12 i8042

(a) The mouse controller, like in IRQ 1. (b) When the mouse moves, it will send interrupts.

### IRQ 14 ata\_piix

(a) Hard drive controller, [Intel PATA/SATA controllers](#). (b) When user read or write files in hardisk, it will send interrupts.

### IRQ 15 ata\_piix

(a) Second hard drive controller. (b) Just like IRQ\_14.

### **IRQ 16 enp0s8**

(a) Ethernet network peripheral, (b) en means ethernet, p0 is bus number 0 and s3 is slot number 8, it is an example of [name scheme](#) for a network device. It will generate interrupts when the system needs to send a package or to deal with a received package.

### **IRQ 18 vmwgfx**

(a) A graphics driver for Linux from VMware, because we use the VM SVGA in our virtual box setting. If we turn to use VBox SVGA, this controller will become vboxvideo. (b) It provides an [acceleration architecture for 2D and 3D](#). When something happens in the GPU it will raise an interrupt.

### **IRQ 19 enp0s3**

Another ethernet interface like the one in IRQ 16, this one is in slot 3.

### **IRQ 20 vboxguest**

(1) This comes from VBox Guest Addition installed in VM, (b) It [improves the usability and performance](#) of guest operating system runs in virtualbox. When the guest needs to synchronize time with the host, it will call an interrupt.

### **IRQ 21 ahci[0000:00:0d.0], snd\_intel8x0**

(1) The sound controller. (2) It is used to manage the sound in Linux system, including generating sound from a device like microphone to application and delivering sound from an application to an output device like earphone.

### **IRQ 22 ohci\_hcd:usb1**

(1) USB interface. (2) [Enables a USB hardware to communicate with a host controller driver in software](#). When a USB device is plugged in or out it will raise an interrupt.