

Lab 5

Members: Fangwen Liao (3439869) | Yijin Wang (3476217)

Exercise 1

The Linux kernel enforces the following policy: a process is not allowed to try to acquire a semaphore while it is already holding a spin lock.

(a) Explain why this policy has to be enforced.

When a process holds a spin lock, it will loop until it gets the resource, it should not wait for a long time. When holding a semaphore, a process can go to sleep, so a process can not go to sleep while holding a spin lock, or when it sleeps it can not release the spin lock.

(b) What can happen when the policy is not enforced? The process holding a spin lock and a semaphore can go to sleep and never release the spin lock, this may cause a deadlock, or it should go to sleep, but spinning and waste CPU cycles. If a process grabs a spinlock and goes to sleep before releasing it. A second process (or an interrupt handler) that wants to grab the spinlock will busy wait. On a uniprocessor machine the second process will lock the CPU not allowing the first process to wake up and release the spinlock so the second process can continue, it is basically a deadlock.

Exercise 2

Describe the problem that occurs when using spinlocks on a single processor system. Consider both the case of

(a) a preemptible kernel and

When a process holding a spinlock is preempted, it doesn't release the lock, if the preempting process needs the resource, it will never get it, thus forming a deadlock. Spinlock automatically disables preemption, which avoids a deadlock caused by interrupts. When data is shared with an interrupt handler, before holding a spinlock we must disable interrupts.

(b) a non-preemptible kernel.

In a non-preemptible kernel, a process will release the resource only when it finishes using it and will not cause a deadlock for the next process.

Exercise 3

Explain what concurrency issues can arise when the `up()` (increment counter) and `down()` (decrement counter) functions of a semaphore are not executed atomically. Illustrate these issues with an example using 2 threads.

When Thread0 acquires the semaphore, the count should be decreased from 1 to zero, but as `down()` is not atomically run, the count stays 1, thus Thread1 can also get the resource, thus causing data inconsistency. When Thread0 finishes using the resource and releases the semaphore, as `up()` is not executed atomically, Thread1 will never be able to get the resource.

Download OS_Lab5.zip and unzip it. Open the kernel directory. It is recommended to make a linked clone of your virtual machine before continuing.

Exercise 4

*Inspect the .sh scripts and download + compile the Linux kernel using these scripts.

(a) Show the output of `uname -a` before installation.

The provided sh files have problems when running, to solve this problem, we copy them and run them inside the virtual machine.

```
fangwenliao@debian:~/Downloads/OS_Lab5/kernel$ uname -a
Linux debian 4.9.0-16-686 #1 SMP Debian 4.9.272-2 (2021-07-19) i686 GNU/Linux
```

(b) Install the four Debian packages that are created once the compilation of the kernel completes successfully.

To run the compile.sh, we first need to apt install time or it will have errors when running the script, or we can just delete the time, since only the make command really matters. After the compile.sh is done, change to root user and use the -i option of dpkg to install the package.

```
root@debian:/home/yijinwang/Documents/OS_Lab5/kernel# dpkg -i linux-headers-
4.19.152_4.19.152-1_i386.deb
Selecting previously unselected package linux-headers-4.19.152.
(Reading database ... 166573 files and directories currently installed.)
Preparing to unpack linux-headers-4.19.152_4.19.152-1_i386.deb ...
Unpacking linux-headers-4.19.152 (4.19.152-1) ...
Setting up linux-headers-4.19.152 (4.19.152-1) ...
```

```
root@debian:/home/yijinwang/Documents/OS_Lab5/kernel# dpkg -i linux-libc-
dev_4.19.152-1_i386.deb
(Reading database ... 190359 files and directories currently installed.)
Preparing to unpack linux-libc-dev_4.19.152-1_i386.deb ...
Unpacking linux-libc-dev (4.19.152-1) over (4.9.290-1) ...
Setting up linux-libc-dev (4.19.152-1) ...
```

```
root@debian:/home/yijinwang/Documents/OS_Lab5/kernel# dpkg -i linux-image-
4.19.152_4.19.152-1_i386.deb
Selecting previously unselected package linux-image-4.19.152.
(Reading database ... 190543 files and directories currently installed.)
Preparing to unpack linux-image-4.19.152_4.19.152-1_i386.deb ...
Unpacking linux-image-4.19.152 (4.19.152-1) ...
Setting up linux-image-4.19.152 (4.19.152-1) ...
update-initramfs: Generating /boot/initrd.img-4.19.152
```

```
Generating grub configuration file ...
Found background image: /usr/share/images/desktop-base/desktop-grub.png
Found linux image: /boot/vmlinuz-4.19.152
Found initrd image: /boot/initrd.img-4.19.152
Found linux image: /boot/vmlinuz-4.9.0-17-686
Found initrd image: /boot/initrd.img-4.9.0-17-686
Found linux image: /boot/vmlinuz-4.9.0-16-686
Found initrd image: /boot/initrd.img-4.9.0-16-686
Found linux image: /boot/vmlinuz-4.9.0-7-686
Found initrd image: /boot/initrd.img-4.9.0-7-686
done
```

```
root@debian:/home/yijinwang/Documents/OS_Lab5/kernel# dpkg -i linux-image-
4.19.152-dbg_4.19.152-1_i386.deb
Selecting previously unselected package linux-image-4.19.152-dbg.
(Reading database ... 190735 files and directories currently installed.)
Preparing to unpack linux-image-4.19.152-dbg_4.19.152-1_i386.deb ...
Unpacking linux-image-4.19.152-dbg (4.19.152-1) ...
Setting up linux-image-4.19.152-dbg (4.19.152-1) ...
```

(c) Reboot into the newly installed kernel. Show the output of **uname a**.

```
fangwenliao@debian:~$ uname -a
Linux debian 4.19.152 #1 SMP Sat Jan 8 08:30:56 CET 2022 i686 GNU/Linux
```

Exercise 5

Implement your own system call that (1) takes one argument (a nice value), (2) prints the current process' PID and nice value, (3) changes the current process' nice value to the nice value that is given as a parameter of the system call, and (4) finally returns the process's PID to user space.

(a) Implement your system call in `sys.c` using the appropriate macro,

our code is as below, the `task_tgid_vnr`, `find_task_by_vpid`, `task_nice` and `set_user_nice` are all used in `sys.c` before, it is fine to use them here, we put them between line 2579 and line 2591, note it can not be put inside `#ifdef`, otherwise it will get errors while compiling.

```
SYSCALL_DEFINE1(newcall, int, nice_value)
{
    int pid;
    struct task_struct *p;
    int ni;

    pid = task_tgid_vnr(current);
    p = find_task_by_vpid(pid);
    ni = task_nice(p);
```

```

    printk("PID: %d, Nice Value:%d.\n", pid, ni);
    set_user_nice(p, nice_value);
    return pid;
}

```

(b) Add an entry for the new system call to the `syscall_32.tbl` for your appropriate architecture,

we enter the following code in `syscall_32.tbl`

```

387      i386      newcall          sys_newcall
__ia32_sys_newcall

```

(c) Compile, install, and boot the kernel containing your system call,

repeat the process in 4

(d) Implement a C application (in user space) that calls your newly implemented system call using the `syscall` function,

our code is like below:

```

#include <stdio.h>

int main(void)
{
    int pid;

    pid = syscall(387, 11);
    printf("Current PID: %d", pid);
    while(1)
    {

    }

    return pid;
}

```

(e) Run your application and show that the system call is made by inspecting the strace.

```

fangwenliao@debian:~/Downloads/kernel$ gcc test.c
test.c: In function 'main':
test.c:7:8: warning: implicit declaration of function 'syscall' [-Wimplicit-
function-declaration]
    pid = syscall(387, 11);
    ^
fangwenliao@debian:~/Downloads/kernel$ strace ./a.out
execve("./a.out", [ "./a.out" ], [ /* 48 vars */ ]) = 0
brk(NULL)                               = 0x10cf000

```

```

access("/etc/ld.so.nohwcap", F_OK)      = -1 ENOENT (No such file or directory)
mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0xb7ed9000
access("/etc/ld.so.preload", R_OK)      = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=100353, ...}) = 0
mmap2(NULL, 100353, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb7ec0000
close(3)                                = 0
access("/etc/ld.so.nohwcap", F_OK)      = -1 ENOENT (No such file or directory)
open("/lib/i386-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\1\1\1\3\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\204\1\0004\0\0\0"...
512) = 512
fstat64(3, {st_mode=S_IFREG|0755, st_size=1791908, ...}) = 0
mmap2(NULL, 1800700, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) =
0xb7d08000
mprotect(0xb7eb9000, 4096, PROT_NONE)   = 0
mmap2(0xb7eba000, 12288, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1b1000) = 0xb7eba000
mmap2(0xb7ebd000, 10748, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7ebd000
close(3)                                = 0
set_thread_area({entry_number:-1, base_addr:0xb7eda100, limit:1048575,
seg_32bit:1, contents:0, read_exec_only:0, limit_in_pages:1, seg_not_present:0,
useable:1}) = 0 (entry_number:6)
mprotect(0xb7eba000, 8192, PROT_READ)   = 0
mprotect(0x453000, 4096, PROT_READ)     = 0
mprotect(0xb7f03000, 4096, PROT_READ)   = 0
munmap(0xb7ec0000, 100353)              = 0
syscall_387(0xb, 0x1, 0x4525e7, 0x1, 0xbfefcc24, 0xbfefcc2c) = 0x6ab
fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
brk(NULL)                               = 0x10cf000
brk(0x10f0000)                          = 0x10f0000
^Z
[1]+  Stopped                  strace ./a.out

```

(f) Show that the nice value of the process has changed. TIPS: use the current macro, and `set_user_nice` function

we use the `ps` command to check the nice value of `a.out`, as is shown the nice value has changed to the value we wanted.

```

fangwenliao@debian:~/Downloads/kernel$ ps -al

```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	117	761	757	0	80	0	-	19524	-	tty1	00:00:00	gnome-session-
0	S	117	769	761	1	80	0	-	226673	-	tty1	00:00:03	gnome-shell
0	S	117	791	769	0	80	0	-	36889	-	tty1	00:00:00	Xwayland
0	S	117	827	761	0	80	0	-	129986	-	tty1	00:00:00	gnome-settings
4	S	1000	1146	1144	1	80	0	-	45691	epoll_	tty2	00:00:02	Xorg
0	S	1000	1160	1144	0	80	0	-	19557	poll_s	tty2	00:00:00	gnome-session-
0	S	1000	1290	1160	4	80	0	-	237140	poll_s	tty2	00:00:09	gnome-shell
0	S	1000	1392	1160	0	80	0	-	105763	poll_s	tty2	00:00:00	gnome-settings
0	S	1000	1412	1160	6	99	19	-	49561	poll_s	tty2	00:00:12	tracker-extrac

0	S	1000	1415	1160	0	80	0	-	40516	poll_s	tty2	00:00:00	gnome-software
0	S	1000	1416	1160	53	99	19	-	34674	poll_s	tty2	00:01:43	tracker-miner-
0	S	1000	1417	1160	0	99	-	-	15567	poll_s	tty2	00:00:00	tracker-miner-
0	S	1000	1420	1	0	80	0	-	18219	poll_s	tty2	00:00:00	gsd-printer
0	S	1000	1421	1160	0	80	0	-	59929	poll_s	tty2	00:00:00	evolution-alar
0	S	1000	1424	1160	0	99	-	-	21444	poll_s	tty2	00:00:00	tracker-miner-
0	T	1000	1705	1630	0	80	0	-	709	signal	pts/0	00:00:00	strace
0	t	1000	1707	1705	42	91	11	-	553	ptrace	pts/0	00:00:07	a.out
0	R	1000	1708	1630	0	80	0	-	1851	-	pts/0	00:00:00	ps